

# Intelligent Robotics Project

- Dhruv Dhamani
- ddhamani@uncc.edu

```
import torch
import torch.nn as nn
import torchvision
%pylab inline

import numpy as np
import pandas as pd
import os
from tqdm import tqdm

from quaternion_layers import *

device = "cuda:0" if torch.cuda.is_available() else "cpu"
```

Populating the interactive namespace from numpy and matplotlib

## The Robot

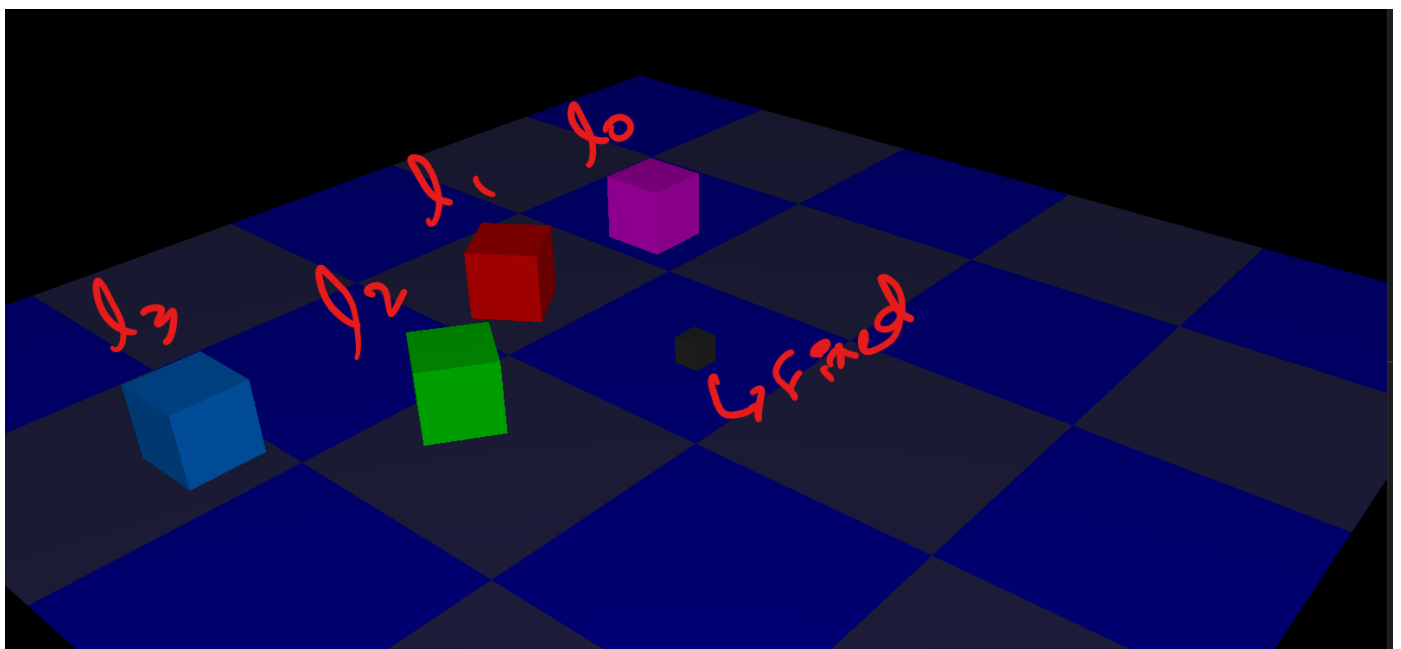
First let's take a look at the robot arm we'll be working with -

```
fixed [⌵]
torso_linear [↑+Z]
shoulder_yaw [⌵+Z]
elbow_yaw [⌵+Z]
wrist_yaw [⌵+Z]
```

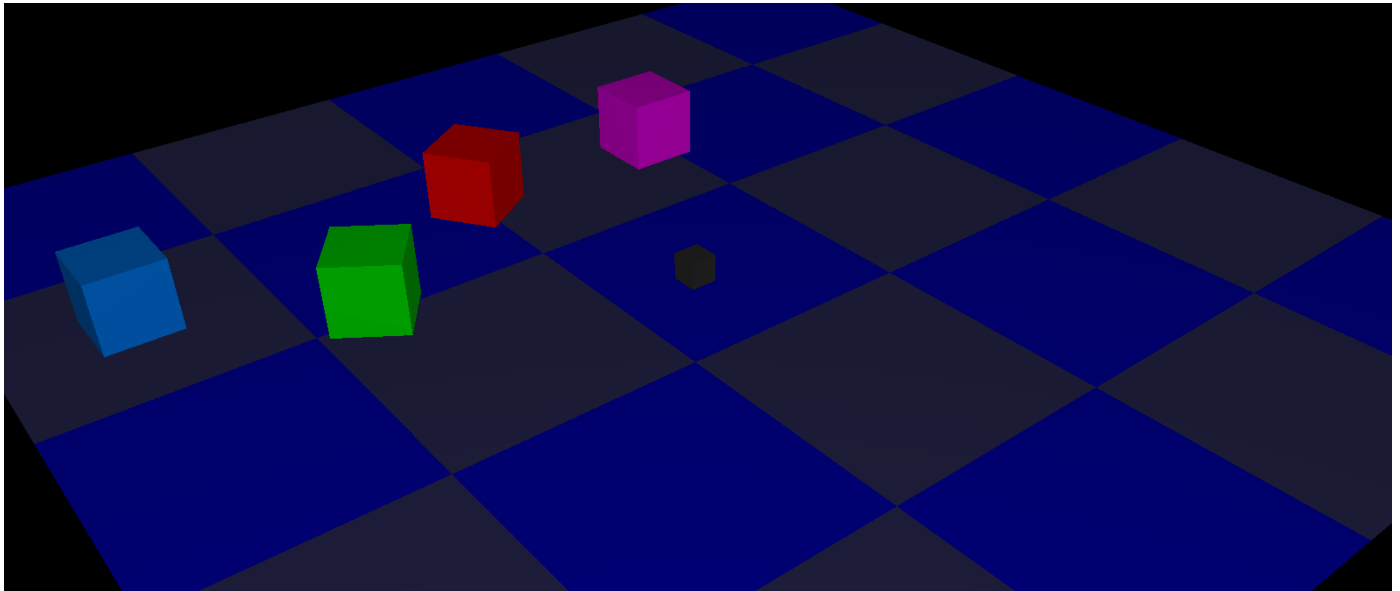
- Fixed joint shown by ⌵
- Linear joint along the Z-axis shown by ↑+Z
- Rotational joint along the Z-axis shown by ⌵+Z

I had trouble coming up with a good configuration for a robot arm with enough "reach" to be able to produce a nice and varied dataset, and this was the best I could come up with at the time.

This is what it looks like -



And this is how it can move -



## The Data

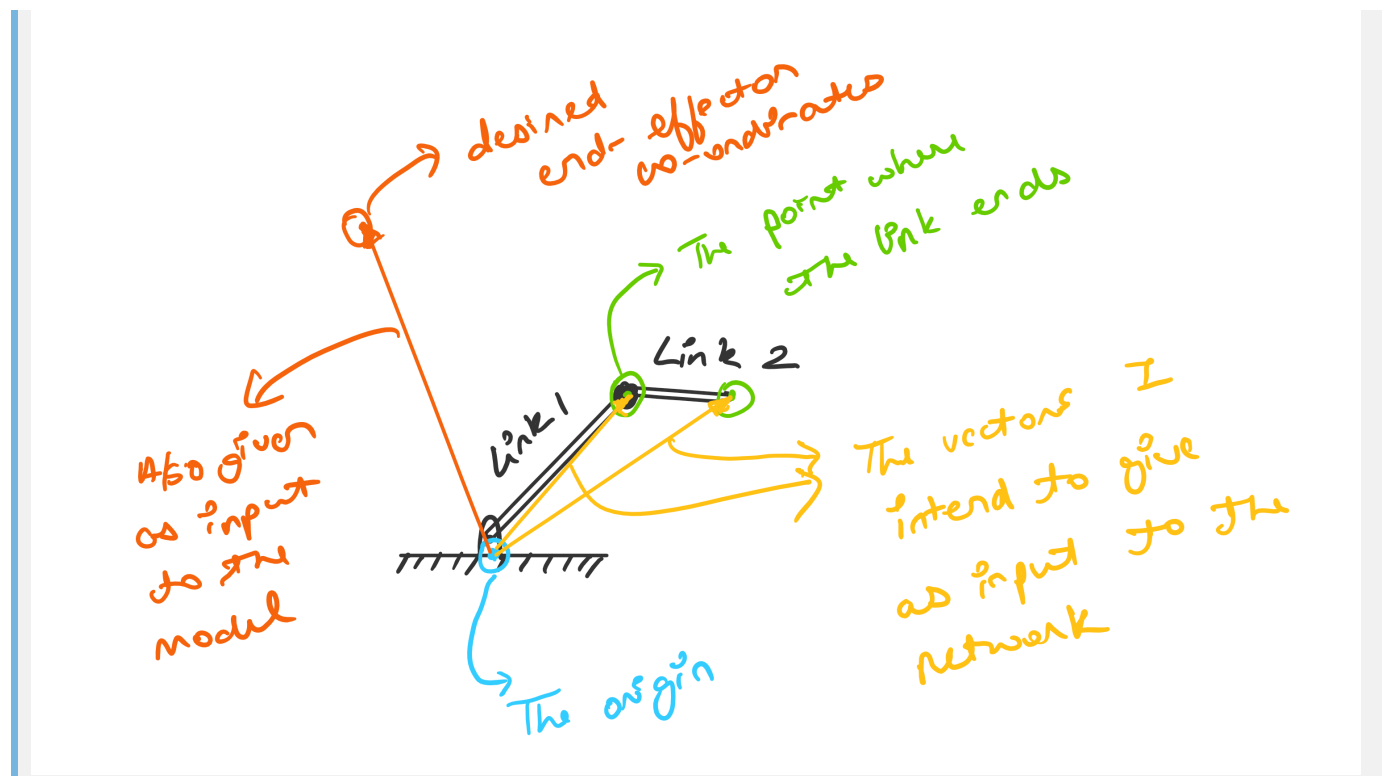
To quote my project proposal, I had said -

A 2R, 3R, and 4R robotic arm would be simulated using ROS and Gazebo. The user would input either via console or through a web interface the desired co-ordinates for the end-effector.

**The desired end-effector co-ordinates would be fed into a neural network as in the form of a quaternion that describes the transformation required to go from the origin to the desired end-effector co-ordinates, along with the quaternions describing the current orientation of each link in the robotic arm.**

As such the data we need to gather is -

- Desired end-effector co-ordinates
- A vector for each link that goes from the origin to the end of that link as shown below (as a quaternion)



- A quaternion for the initial orientation of the link

These would be the input to the network; for training, we also need the corresponding outputs, which would be -

- The translation needed to be made for each link, as a quaternion.
- The rotation needed to be made for each link, as a quaternion.

I computed all of this data by writing a script in Rust, using a kinematics library called [k](#). The data computed was written to a [JSON](#) file, as an array of [JSON](#) objects, each representing one input-output pair. One such input-output pair is detailed below -

```
{
  "l0": [-0.0008749962, 1.0989358, 0.10087502],
  "l1": [-0.00087475777, 1.098697, 0.4008749],
  "l2": [0.047727853, 1.0989714, 0.10483825],
  "l3": [-0.24934354, 1.0987016, 0.14665347],
  "l0_rot": [-0.49960187, -0.49999982, -0.49999982, 0.5003981],
  "l1_rot": [-0.53895015, 0.45731235, 0.45810595, 0.5390149],
  "l2_rot": [-0.04910004, -0.7051177, -0.7056401, 0.049701005],
  "l3_rot": [-0.49960187, -0.49999973, -0.49999982, 0.50039816],
  "l0_final": [-0.000079125166, 0.09939951, 0.100079186],
  "l1_final": [-0.00007888675, 0.09916064, 0.40007907],
  "l2_final": [0.009956747, 0.099407375, 0.100247085],
  "l3_final": [-0.14984122, 0.09948231, -0.15365165],
  "l0_rot_final": [-0.49960187, -0.49999982, -0.49999982, 0.5003981],
  "l1_rot_final": [-0.5082875, 0.49116763, 0.49196374, 0.50830084],
  "l2_rot_final": [0.34192163, -0.6186206, -0.6193856, -0.34170097],
  "l3_rot_final": [-0.49960193, -0.4999997, -0.49999976, 0.50039816],
  "l0_trans": [-0.00015924126, 0.19999987, 0.00015926361],
  "l1_trans": [-0.000079125166, 0.09939951, 0.100079186],
  "l2_trans": [-0.00007888675, 0.09916064, 0.40007907],
  "l3_trans": [0.009956747, 0.099407375, 0.100247085],
  "l0_rot_trans": [-0.49960187, -0.49999982, -0.49999982, 0.5003981],
  "l1_rot_trans": [-0.49960187, -0.49999982, -0.49999982, 0.5003981],
  "l2_rot_trans": [-0.5082875, 0.49116763, 0.49196374, 0.50830084],
  "l3_rot_trans": [0.34192163, -0.6186206, -0.6193856, -0.34170097],
  "a_joint_pos": [0.7990161, 2.978866, -4.409823, 1.4309571],
  "b_joint_pos": [-0.20052081, 3.1081338, -5.6879864, 2.5798528]
}
```

The following piece of code reads the aforementioned `JSON` file and makes a pandas dataframe out of it, which we will then convert to PyTorch tensors and feed them into the network.

```
data = pd.read_json('DATA.json')

data['l0'] = data['l0'].apply(lambda x: np.array([0.0] + x))
data['l1'] = data['l1'].apply(lambda x: np.array([0.0] + x))
data['l2'] = data['l2'].apply(lambda x: np.array([0.0] + x))
data['l3'] = data['l3'].apply(lambda x: np.array([0.0] + x))
data['l0_rot'] = data['l0_rot'].apply(np.array)
data['l1_rot'] = data['l1_rot'].apply(np.array)
data['l2_rot'] = data['l2_rot'].apply(np.array)
data['l3_rot'] = data['l3_rot'].apply(np.array)

data['l0_final'] = data['l0_final'].apply(lambda x: np.array([0.0] + x))
data['l1_final'] = data['l1_final'].apply(lambda x: np.array([0.0] + x))
data['l2_final'] = data['l2_final'].apply(lambda x: np.array([0.0] + x))
data['l3_final'] = data['l3_final'].apply(lambda x: np.array([0.0] + x))
data['l0_rot_final'] = data['l0_rot_final'].apply(np.array)
data['l1_rot_final'] = data['l1_rot_final'].apply(np.array)
data['l2_rot_final'] = data['l2_rot_final'].apply(np.array)
data['l3_rot_final'] = data['l3_rot_final'].apply(np.array)

data['l0_trans'] = data['l0_trans'].apply(lambda x: np.array([0.0] + x))
data['l1_trans'] = data['l1_trans'].apply(lambda x: np.array([0.0] + x))
data['l2_trans'] = data['l2_trans'].apply(lambda x: np.array([0.0] + x))
data['l3_trans'] = data['l3_trans'].apply(lambda x: np.array([0.0] + x))
data['l0_rot_trans'] = data['l0_rot_trans'].apply(np.array)
data['l1_rot_trans'] = data['l1_rot_trans'].apply(np.array)
data['l2_rot_trans'] = data['l2_rot_trans'].apply(np.array)
data['l3_rot_trans'] = data['l3_rot_trans'].apply(np.array)

# data = data[ ['l0', 'l1', 'l2', 'l3', 'l3_final',
#               'l0_rot', 'l1_rot', 'l2_rot', 'l3_rot', 'l3_rot_final',
#               'l0_trans', 'l1_trans', 'l2_trans',
#               'l0_rot_trans', 'l1_rot_trans', 'l2_rot_trans',
#               'l0_trans', 'l1_trans', 'l2_trans',
#               'l0_rot_trans', 'l1_rot_trans', 'l2_rot_trans'] ]

_data = data[ ['l0', 'l1', 'l2', 'l3', 'l3_final',
               'l0_rot', 'l1_rot', 'l2_rot', 'l3_rot',
               'l0_trans', 'l1_trans', 'l2_trans', 'l3_trans',
               'l0_rot_trans', 'l1_rot_trans', 'l2_rot_trans', 'l3_rot_trans'] ]
_data.head()
```

```
.dataframe tbody tr th {
    vertical-align: top;
}

.dataframe thead th {
    text-align: right;
}
```

	I0	I1	I2	I3	I3_final	I0_rot	I1_rot
0	[0.0, -0.00019985437, 0.2510149, 0.10019991]	[0.0, -0.0001996159599999998, 0.25077602, 0.4...	[0.0, 0.034441292000000005, 0.25056633, 0.698193]	[0.0, -0.24929011, 0.25026283, 0.7956424400000...	[0.0, 0.14988443, -0.14983344, -0.103550166000...	[-0.49960187, -0.49999982, -0.49999982, 0.5003...	[-0.52768135, -0.4702704, -0.47022435, 0.5...
1	[0.0, 0.000119552016, -0.15011469, 0.09988052]	[0.0, 0.000119790435, -0.15035355, 0.399880399...	[0.0, 0.21400562, -0.15001574, 0.18951705]	[0.0, 0.14988443, -0.14983344, -0.103550166000...	[0.0, 0.24937618, -0.15022174, -0.10424779]	[-0.49960187, -0.49999982, -0.49999982, 0.5003...	[-0.65424514, 0.26751652, 0.26825088, 0.65455294]
2	[0.0, 0.000119924545, -0.15058279, 0.09988013]	[0.0, 0.000120162964, -0.15082166, 0.39988002]	[0.0, 0.21847203, -0.15048397, 0.1941560700000...	[0.0, 0.24937618, -0.15022174, -0.10424779]	[0.0, -0.24936235, -0.050322868, -0.10379206]	[-0.49960187, -0.49999982, -0.49999982, 0.5003...	[-0.6570787, 0.2604791999999, 0.26121026, ...]
3	[0.0, 4.0054319999999995e-05, -0.0502865, 0.09...	[0.0, 4.029274e-05, -0.050525367, 0.39995986]	[0.0, -0.030634344, -0.05031219, 0.10153228]	[0.0, -0.24936235, -0.050322868, -0.10379206]	[0.0, -0.15061736, -0.05022552, -0.10425120600...	[-0.49960187, -0.49999982, -0.49999982, 0.5003...	[-0.473769, 0.52454245, 0.52533764, 0.47372824]
4	[0.0, 4.0039419999999996e-05, -0.0502681730000...	[0.0, 4.027784e-05, -0.05050704, 0.3999599]	[0.0, 0.06279153, -0.050223485000000005, 0.106...	[0.0, -0.15061736, -0.05022552, -0.10425120600...	[0.0, -0.05064021, -0.05022476, -0.10428929]	[-0.49960187, -0.49999982, -0.49999982, 0.5003...	[-0.54976964000, 0.4442468, 0.445, 0...

The columns -

- I0
- I1
- I2
- I3
- I0\_rot
- I1\_rot
- I2\_rot
- I3\_rot
- I3\_final

will be given to the network as input, and we'll expect the following as output -

- I0\_trans
- I1\_trans
- I2\_trans
- I3\_trans
- I0\_rot\_trans
- I1\_rot\_trans
- I2\_rot\_trans
- I3\_rot\_trans

```
_data = np.array([np.hstack(list(x)) for x in list(_data.values)])
_data.shape
```

```
(136446, 68)
```

We have 136,446 input-output samples of data.

```
def train(_data, model):
    loss = torch.nn.MSELoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=0.01, amsgrad=True)
```

## Experiment - 1

- 
- The graph displays the function  $f(x) = \frac{1}{1+e^x}$ . The x-axis is labeled from 0 to 100 in increments of 20. The y-axis is labeled from 0.0 to 1.0 in increments of 0.2. The curve begins at the point (0, 1.0) and decreases sharply, reaching a value of approximately 0.05 at  $x=10$ , and then continues to decrease very slowly, approaching 0 as  $x$  increases towards 100.

loss\_curve

```
[0.986152123002445,  
0.2368584047345554,  
0.05968216587515438,  
0.03567931198460214,  
0.02652234474525732,  
0.021952509660931194,  
0.019351086360128486,  
0.01769569784621982,  
0.01656747828511631,  
0.015746465281528586,  
0.015109013574307455,  
0.01458377024049268,  
0.014166720375856933,  
0.013820498962612712,  
0.013551476521088797,  
0.013257118188502157,  
0.013060236431877403,  
0.012852968675467898,  
0.012670425013365114,  
0.012520903611884397,  
0.012369291172088945,  
0.012257742914645112,  
0.012124518973424155,  
0.012018682325587553,  
0.011947989189887749,  
0.01183365737361943,  
0.011756880318417269,  
0.011689543066655888,  
0.011624884490361986,  
0.011556160143193077,  
0.011481914231005837,  
0.011442565249607843,  
0.011397028122754657,  
0.011335667113170904,  
0.011283821125021753,  
0.011273231987348375,  
0.011200207422542222,  
0.011174608706770575,  
0.011120990055668004,  
0.011117238277459847,  
0.011093359741875353,  
0.011038905858774395,  
0.01100745375322945,  
0.010994048871319084,  
0.010950398877920473,  
0.01094581778435146,  
0.010915991274968666,  
0.010905670056886533,  
0.010877651326796588,  
0.010852692730943947,  
0.01084841916556744,  
0.01082376141429824,  
0.010805864078814493,  
0.01077498895499636,  
0.01076501076493193,  
0.010752176592016922,  
0.010734438512693434,  
0.010737747282666318,  
0.010716543532907963,  
0.010701483937309068,  
0.010686827768736026,  
0.010681208606590243,  
0.010673078270081212,  
0.010677565765731475,  
0.010653222308439366,  
0.010644585058531341,  
0.010642239193925086,  
0.010632648543619058,  
0.010603149120202836,  
0.0106283373418538,  
0.010607976658160196,  
0.010597313261207412,
```

## Experiment - 2

[0.6961357540944043,  
0.1277247665121275,  
0.017848176580360708,  
0.011656587779083672,  
0.010295450906543172,  
0.009626669097034371,  
0.009234811864135897,  
0.008998205063535887,  
0.00883806957041516,  
0.008757768664509058,  
0.008674929233487038,  
0.008630803868393688,  
0.008597744053558391,  
0.008568069705849184,  
0.008559234905987978,  
0.008539102040231228,  
0.008496531394912916,  
0.008482503649943015,  
0.008468611850677168,  
0.0084363707642564,  
0.008392463205382228,  
0.008404746679041316,  
0.008370248819975293,  
0.008347262231194797,  
0.008327387437662658,  
0.008317987160647617,  
0.008284992100123097,  
0.008276611114578211,  
0.00825809468241299,  
0.008245306870187907,  
0.008220543783596334,  
0.008213697727221777,  
0.008191911298233797,  
0.0081790333668537,  
0.008172581410583328,  
0.008135043563978636,  
0.008130770245128694,  
0.008112591540660052,  
0.008100147701471168,  
0.008082948652479579,  
0.008088587660013753,  
0.008062657653628028,  
0.008035115883959568,  
0.008031446212793098,  
0.008037235108478105,  
0.008026150589370552,  
0.00801415227846626,  
0.007990884526139674,  
0.00799490199150408,  
0.007992539335699642,  
0.007958530818166979,  
0.00795279792509973,  
0.007950917666997103,  
0.007945361560867989,  
0.007947011344025241,  
0.00793238149901085,  
0.007921579198035248,  
0.007912321564029245,  
0.007906501715564552,  
0.007897652083021752,  
0.007892324561801027,  
0.007895895636037868,  
0.007859488127424437,  
0.007887378471958287,  
0.007855288862415096,  
0.007853892125079738,  
0.00786541765281821,  
0.007856967461788478,  
0.007841414046090315,  
0.00786130509667975,  
0.007851362680358923,  
0.007843551061609211,  
0.007845936356769764,  
0.007817156284170993,  
0.007817397934987265,  
0.0078098784776075795,



```

0.007817657376803896,
0.007818766924388269,
0.007803790079539313,
0.0078092080349212185,
0.00779493443448754,
0.007809588137794943,
0.007785964962642859,
0.00780522812377004,
0.007782130721299087,
0.007789943901383702,
0.007785777301144074,
0.007784320000449524,
0.007771506111192352,
0.007770315483760308,
0.007766653304262196,
0.007759088129900834,
0.007754348382792052,
0.0077764961649389825,
0.007764776524923304,
0.007774298490189454,
0.0077714372344095915,
0.007742976400015109,
0.007762465434258475,
0.007777250084259054]

```

## Experiment - 3

1. Quaternion Linear Layer (9 quaternions input -> 10 quaternions output)
2. ELU activation.
3. Quaternion Linear Layer (10 quaternions input -> 8 quaternions output)

```

model = nn.Sequential(
    QuaternionLinearAutograd(36, 40),
    nn.ELU(),
    QuaternionLinearAutograd(40, 32)
)
loss_curve, _ = train(_data, model)
plot(loss_curve)

```

```

Sequential(
  (0): QuaternionLinearAutograd(in_features=9, out_features=10, bias=True, init_criterion=glorot, weight_init=quaternion, seed=16)
  (1): ELU(alpha=1.0)
  (2): QuaternionLinearAutograd(in_features=10, out_features=8, bias=True, init_criterion=glorot, weight_init=quaternion, seed=812)
)

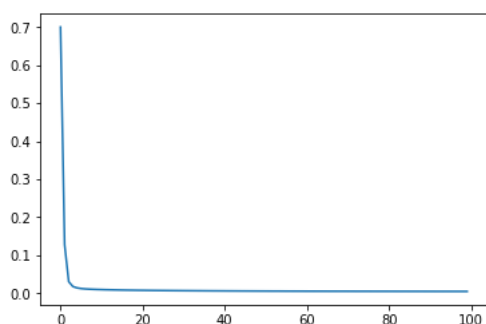
```

```

100%|
100/100 [04:29<00:00, 2.87s/it]

```

```
[<matplotlib.lines.Line2D at 0x21dae94cef0>]
```

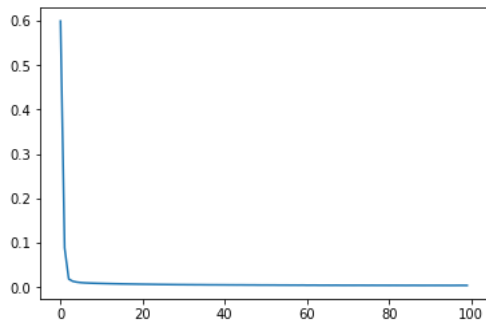


loss\_curve

```
[0.6998992404517006,  
0.12822098953320699,  
0.030986769398783937,  
0.01827019696836086,  
0.014580334180637318,  
0.012843650464406785,  
0.011851914528319064,  
0.011207187000442953,  
0.010748690405093572,  
0.010343435384771404,  
0.010038840200971155,  
0.009734360775088562,  
0.00948428047601791,  
0.009251174763502443,  
0.009043739971649997,  
0.008838466094697224,  
0.008660778074580081,  
0.008516683517133488,  
0.008376131806632174,  
0.008216055808588862,  
0.00812964845339165,  
0.007996211362564388,  
0.007883172879433808,  
0.007772794212488567,  
0.007692442727549111,  
0.007586487500435289,  
0.007525661960244179,  
0.0074337767163182005,  
0.007350209799102124,  
0.007274761617950657,  
0.007202370900332052,  
0.007133643293534131,  
0.007070556652786977,  
0.007001456818269456,  
0.006932493312941755,  
0.0068684721447746545,  
0.006796826457823901,  
0.006731509751475909,  
0.006678657592548167,  
0.006636147674940088,  
0.006564837962608127,  
0.006511765997856855,  
0.006461161298348624,  
0.006427806076210211,  
0.006354283459265442,  
0.006323671954519609,  
0.006287563633283272,  
0.006245004684280823,  
0.006203523631591131,  
0.006162540192770607,  
0.006131131777211148,  
0.006080326966612655,  
0.0060610603404176585,  
0.006054841433925664,  
0.006011319059111616,  
0.00595988637274679,  
0.005958708249689902,  
0.0059225247241556644,  
0.005903343820725293,  
0.005871022161205902,  
0.005834729095701785,  
0.0058238756541601,  
0.005797940361149171,  
0.005795579458422521,  
0.005765828988788759,  
0.005737587220638114,  
0.005712798741810462,  
0.005703149596229196,  
0.0056929219212821305,  
0.005692791344378801,  
0.005666661898002905,  
0.0056573650662732474,
```

## Experiment - 4

- ```
[<matplotlib.lines.Line2D at 0x21dae9ef8d0>]
```

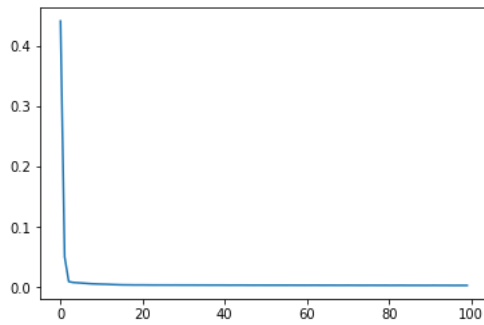


loss\_curve

```
[0.5994927848086637,  
 0.08886091427548844,  
 0.018699251033146593,  
 0.013594621710260124,  
 0.011631623677471104,  
 0.010588998796747011,  
 0.009983046254252686,  
 0.009584668256780681,  
 0.00924441300551681,  
 0.008965311626739362,  
 0.008740470114657107,  
 0.008489847087356098,  
 0.008282916729941088,  
 0.008066133582307136,  
 0.007892180590287727,  
 0.007713019272641224,  
 0.007554194780395311,  
 0.0074010052142993495,  
 0.007257164387470659,  
 0.007110567270394634,  
 0.006979883497800021,  
 0.006861647974480601,  
 0.006751525010365774,  
 0.006613005870295798,  
 0.006470081355313168,  
 0.00641014811801998,  
 0.006290620987248772,  
 0.006213603842565242,  
 0.006104922565795919,  
 0.006024379652979619,  
 0.00600840404684491,  
 0.005893863806062762,  
 0.005881903857430991,  
 0.0057892020353499584,  
 0.005789527450414265,  
 0.005709070849287159,  
 0.0056217447854578495,  
 0.0055879187112783685,  
 0.0055845681824447475,  
 0.005522696315036977,  
 0.005499558797215714,  
 0.005417376249919043,  
 0.005427693044218947,  
 0.0053521492107606984,  
 0.005393465611097568,  
 0.005306020447546069,  
 0.005254544440985602,  
 0.005265700398012996,  
 0.005153646185884581,  
 0.005108028787243015,  
 0.00513393894824035,  
 0.005076153340804226,  
 0.005097979118170983,  
 0.004969004130757907,  
 0.0049698457450551145,  
 0.004960205049856621,  
 0.00487235067960094,  
 0.004908610162708689,
```

## Experiment - 5

- ```
[<matplotlib.lines.Line2D at 0x21dc3a139e8>]
```



loss\_curve

```
[0.4406554935609593,  
0.05094682819703046,  
0.009803676060126984,  
0.008327381918206811,  
0.007860258413369165,  
0.007413569577586125,  
0.006870988701634547,  
0.00634442810790942,  
0.005963136359830113,  
0.005706874338690849,  
0.005498124943936572,  
0.005329106304356281,  
0.005073142977540984,  
0.004796549664152896,  
0.004557841840912314,  
0.004377546480528134,  
0.004294636494973127,  
0.004243447616532007,  
0.004138304810861454,  
0.00414498956353568,  
0.004157040672212401,  
0.004085841465412694,  
0.00405130788738675,  
0.004018921131158576,  
0.004025966860353947,  
0.004020402837982949,  
0.004006334847551496,  
0.003973330542280832,  
0.004013155477450174,  
0.003969973350382026,  
0.003966285299290629,  
0.00394236679901095,  
0.003959921065389234,  
0.003920264863956939,  
0.003923632293556105,  
0.003929552889209898,  
0.003909661626333699,  
0.003908213618321016,  
0.0039185546954874605,  
0.0039028870626626644,  
0.0038733788738575053,  
0.0038330212084795624,  
0.003842621644520584,  
0.0038632446667179465,  
0.0038537685690885957,  
0.0038412495392977316,  
0.00379936904509497,  
0.003786977380514145,  
0.0037999787395272184,  
0.003791563390918514,  
0.0037817786968148805,  
0.003773194580229328,  
0.003782298231004354,  
0.003770866383360151,  
0.0037563663619734787,  
0.0037530595916049447,
```

```
0.0037892624831703655,  
0.0037531758699675694,  
0.003713937209206907,  
0.0037438271717880577,  
0.003697621834683506,  
0.0037375407581053234,  
0.0037643533158937797,  
0.0037449329256025307,  
0.003669330514693523,  
0.0036873071769471556,  
0.003712779321872136,  
0.0036501581506694063,  
0.003646914443165502,  
0.0036369644666967146,  
0.0037011085352038637,  
0.0036410300458288368,  
0.003623508019646739,  
0.003607176004580277,  
0.0035727926091674495,  
0.003634456866968642,  
0.0036346051458488494,  
0.003565456656574765,  
0.0035428654416190353,  
0.0036139690508956417,  
0.0035359286703169346,  
0.0035447543281513978,  
0.003573565504184979,  
0.0036046415197608225,  
0.0035341630744583465,  
0.0034985918455811985,  
0.0035310751443509668,  
0.0035163988701670487,  
0.0034818547942182597,  
0.0034790953949970357,  
0.003488030229859492,  
0.003555999233332627,  
0.0034880095284760874,  
0.0034853933269486707,  
0.003494919222943923,  
0.0034528779523337588,  
0.0033922527629115127,  
0.003448927747633527,  
0.003425896626568454,  
0.0034207105294198673]
```